

Experiment 1: Working with Python packages-Numpy, Scipy, Scikit-Learn, Matplotlib

Aim: To explore Python packages such as NumPy, SciPy, Pandas, Scikit-learn, and Matplotlib, and apply machine learning workflows on datasets from UCI and Kaggle repositories.

Libraries used:

- Pandas - for data handling
- numpy - for numerical operations
- matplotlib.pyplot - for visualization
- sklearn - for model building and evaluation
- Scipy - provides a large collection of functions for advanced mathematical, scientific, and engineering computations

Exploring Python Libraries

NumPy

- **Key Functions:** array, reshape, arange, linspace, zeros, ones, sum, mean, std, dot, random
- **Operations:** Vectorized computation, matrix algebra, broadcasting

Pandas

- **Common Functions:** DataFrame, read_csv, head(), info(), describe(), groupby(), merge(), pivot_table()
- **Data Cleaning:** dropna(), fillna(), astype(), map(), replace()

SciPy

- **Purpose:** Scientific computation and mathematical operations
- **Key Modules:** scipy.stats (statistical tests), scipy.optimize (optimization), scipy.integrate, scipy.spatial, scipy.signal

Scikit-Learn

- **Tasks:** Classification, regression, clustering, model evaluation
- **Key Modules:** datasets, model_selection, preprocessing, metrics, linear_model, tree, svm

Matplotlib

- **Purpose:** Data visualization and graphical representation
- **Common Functions:** plot(), scatter(), bar(), hist(), imshow(), subplot(), title(), xlabel(), ylabel()

Objective: The objective of this experiment is to explore and understand the functionality of essential Python libraries used in data analysis and machine learning, namely **NumPy**, **Pandas**, **SciPy**, **Scikit-learn**, and **Matplotlib**.

The experiment also aims to:

- Understand the core operations such as array manipulations, data preprocessing, scientific computing, model building, and data visualization.
- Apply machine learning workflows to real-world datasets obtained from public repositories like the UCI Machine Learning Repository and Kaggle.
- Identify suitable machine learning tasks and algorithms for various datasets such as loan prediction, handwritten digit recognition, email spam classification, diabetes prediction, and the Iris dataset.
- Perform data loading, exploratory data analysis, preprocessing, feature selection, model training, and performance evaluation.

ML WORKFLOW STEPS

- Load dataset using Pandas (read_csv)
- EDA: describe(), info(), value_counts(), plotting
- Preprocessing: LabelEncoder, StandardScaler, handling NaN
- Feature Selection: SelectKBest, f_classif, chi2
- Split Data: train_test_split()
- Model Selection & Training
- Performance Evaluation: Accuracy, confusion matrix, classification report

CODE:

```
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
```

-----IRIS DATASET-----

```
# Load dataset
iris = sns.load_dataset('iris')

# EDA
print(iris.head())
sns.pairplot(iris, hue='species')
plt.show()

# Preprocessing
X = iris.drop('species', axis=1)
y = iris['species']

# Feature selection
selector = SelectKBest(score_func=f_classif, k='all')
X_new = selector.fit_transform(X, y)

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(
    X_new, y, test_size=0.2, random_state=42)

# Model
model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)

# Evaluation
y_pred = model.predict(X_test)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

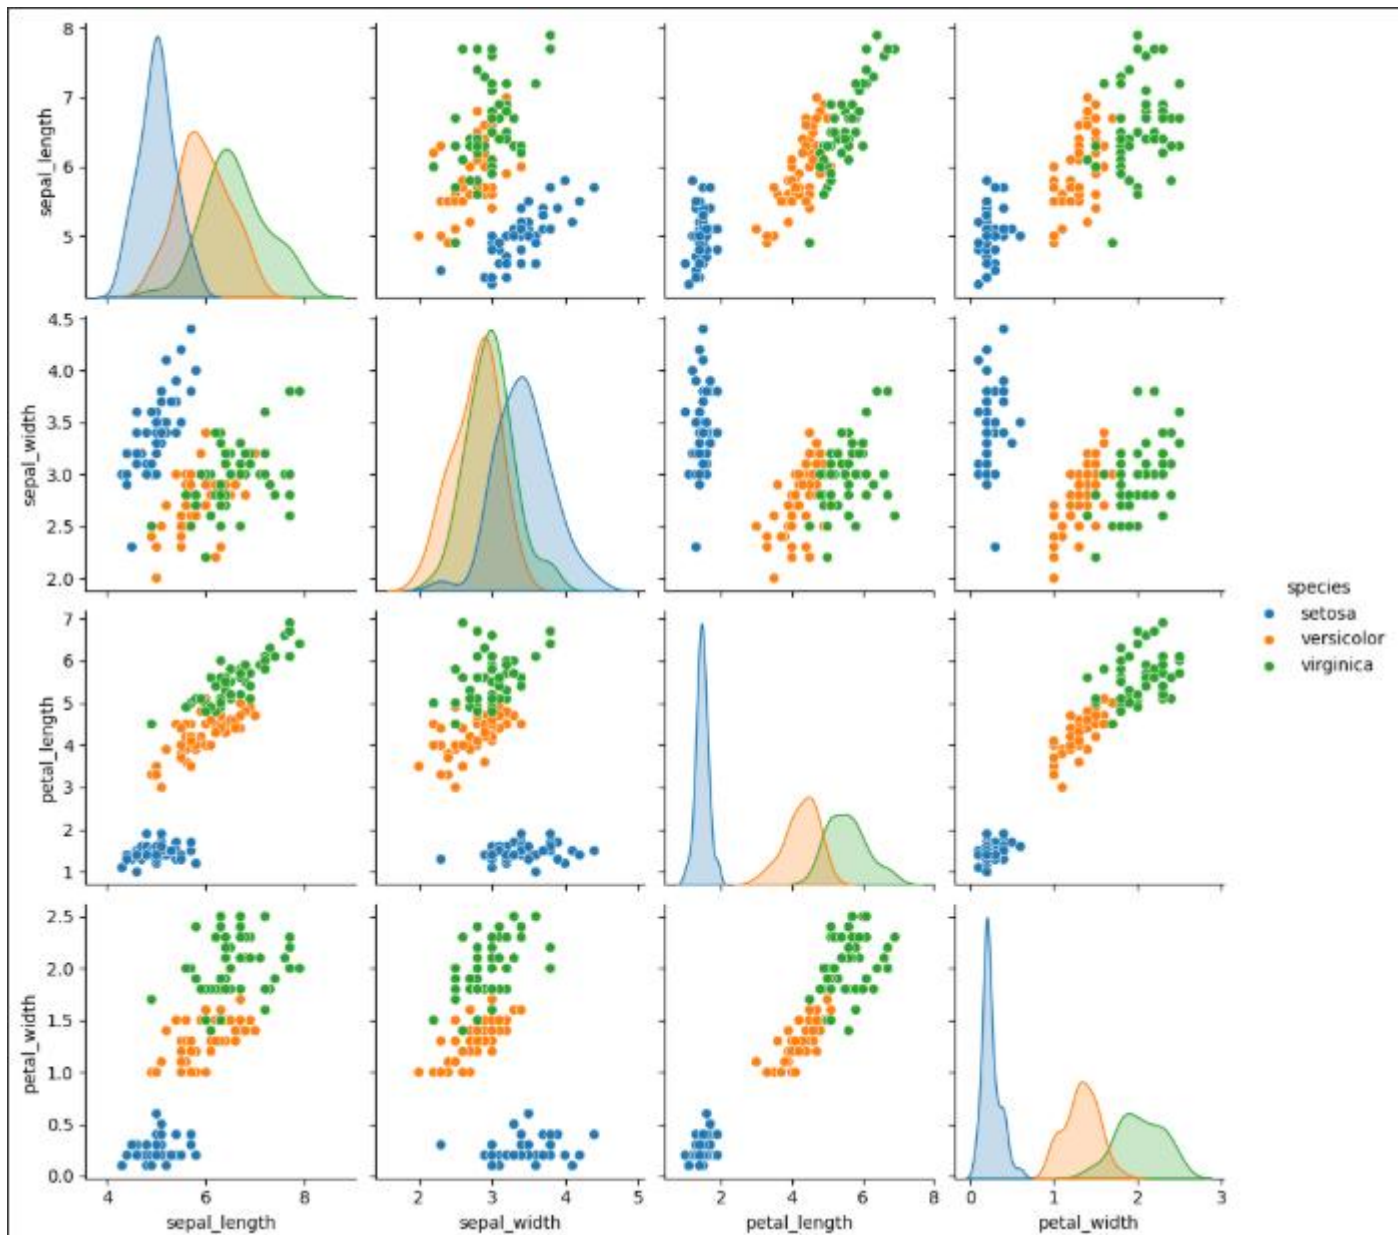
OUTPUT

Date: 23-07-2025

Experiment: 1

Name: Pranavah Varun

Roll No: 3122237001039



-----LOAN AMOUNT DATASET-----

```
# 1. Import libraries
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from google.colab import drive
```

```
# 2. Load dataset from Google Drive
```

```
drive.mount('/content/drive')
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/ML LAB SEM 5/train.csv')

# 3. Drop identifier columns and handle missing values
df.drop(columns=["Customer ID", "Name", "Property ID"], inplace=True)
df.dropna(inplace=True)

# 4. Define features and target
X = df.drop(columns=["Loan Sanction Amount (USD)"])
y = df["Loan Sanction Amount (USD)"]

# 5. Encode categoricals and scale numeric features
X = pd.get_dummies(X, drop_first=True)
X = StandardScaler().fit_transform(X)

# 6. Train-test split and model training
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)


# 7. Model evaluation
y_pred = model.predict(X_test)
print("RMSE:", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))

# 8. Plotting Actual vs Predicted values
plt.figure(figsize=(8,6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.6, color='royalblue')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual Loan Sanction Amount (USD)")
plt.ylabel("Predicted Loan Sanction Amount (USD)")
plt.title("Actual vs Predicted Loan Amounts")
plt.grid(True)
plt.tight_layout()
plt.show()
```

OUTPUT

RMSE: 1195267145.5071688

R² Score: 0.47512320259332885



-----DIABETES DATASET-----

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score

# Load dataset
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/
ML LAB SEM 5/diabetes.csv') # PIMA Indian dataset

# EDA
print(df.info())
print(df.describe())

# Preprocessing
X = df.drop('Outcome', axis=1)
y = df['Outcome']
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split
X_train, X_test, y_train, y_test = train_test_split
(X_scaled, y, test_size=0.25, random_state=42)

# Model
```

```
clf = RandomForestClassifier()
clf.fit(X_train, y_train)

# Evaluation
y_pred = clf.predict(X_test)
print(accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["No Diabetes", "Diabetes"],
yticklabels=["No Diabetes", "Diabetes"])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.tight_layout()
plt.show()

# Feature Importance
import numpy as np

feature_names = df.columns[:-1] # All columns except 'Outcome'
importances = clf.feature_importances_
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(10,6))
sns.barplot(x=importances[indices], y=feature_names[indices])
plt.title("Feature Importances (Random Forest)")
plt.xlabel("Importance")
plt.ylabel("Features")
plt.tight_layout()
plt.show()

from sklearn.metrics import roc_curve, auc

y_proba = clf.predict_proba(X_test)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
```

```
plt.xlim([-0.01, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Receiver Operating Characteristic")
plt.legend(loc="lower right")
plt.grid(True)
plt.tight_layout()
plt.show()
```

OUTPUT

RangeIndex: 768 entries, 0 to 767

Data columns (total 9 columns):

#	Column	Non-Null	Count	Dtype
0	Pregnancies	768	non-null	int64
1	Glucose	768	non-null	int64
2	BloodPressure	768	non-null	int64
3	SkinThickness	768	non-null	int64
4	Insulin	768	non-null	int64
5	BMI	768	non-null	float64
6	DiabetesPedigreeFunction	768	non-null	float64
7	Age	768	non-null	int64
8	Outcome	768	non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

None

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

0.7395833333333334


Date: 23-07-2025

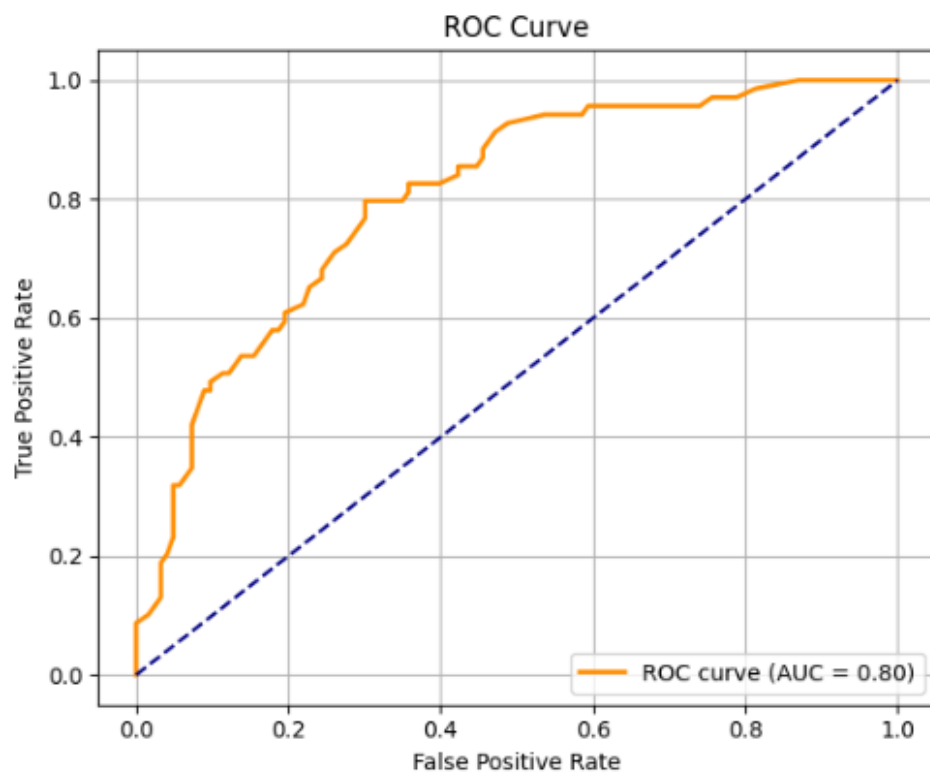

Experiment: 1

Name: Pranavah Varun

Roll No: 3122237001039

	precision	recall	f1-score	support
0	0.80	0.79	0.80	123
1	0.63	0.65	0.64	69
accuracy			0.74	192
macro avg	0.72	0.72	0.72	192
weighted avg	0.74	0.74	0.74	192





-----SPAM DATASET-----

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report
```

```
# Load dataset
df = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/ML LAB SEM 5/
spambase.data", header=None) # UCI SpamBase dataset

# Preprocessing
X = df.iloc[:, :-1]
y = df.iloc[:, -1]

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)


# Model
model = GaussianNB()
model.fit(X_train, y_train)

# Evaluation
y_pred = model.predict(X_test)
print(accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

OUTPUT

0.8067318132464713

	precision	recall	f1-score	support
0	0.94	0.72	0.81	538
1	0.70	0.93	0.80	383
accuracy			0.81	921
macro avg	0.82	0.82	0.81	921
weighted avg	0.84	0.81	0.81	921



-----MNIST DATASET-----

```
# Load dataset
```

```
digits = load_digits()
```

```
X = digits.data
```

```
y = digits.target
```

```
# Show a sample digit
```

```
plt.imshow(digits.images[0], cmap='gray')
```

```
plt.title(f'Label: {digits.target[0]}')
```

```
plt.show()
```

```
# Split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Model
```

```
model = LogisticRegression(max_iter=3000)
```

```
model.fit(X_train, y_train)
```

```
# Evaluation
```


```
y_pred = model.predict(X_test)
```

```
print(accuracy_score(y_test, y_pred))
```

```
print(classification_report(y_test, y_pred))
```

OUTPUT

	precision	recall	f1-score	support
0	1.00	1.00	1.00	53
1	0.96	0.94	0.95	50
2	0.96	1.00	0.98	47
3	0.98	0.96	0.97	54
4	1.00	0.97	0.98	60
5	0.94	0.95	0.95	66
6	0.96	0.98	0.97	53
7	1.00	0.96	0.98	55
8	0.91	0.98	0.94	43
9	0.97	0.95	0.96	59
accuracy			0.97	540
macro avg	0.97	0.97	0.97	540
weighted avg	0.97	0.97	0.97	540



Results and Discussions:

Dataset	ML Task	ML Type	Dataset Source
Loan Amount Prediction	Regression	Supervised	Kaggle / UCI Repository
Handwritten Character Recognition	Classification	Supervised	MNIST Dataset
Email Spam Classification	Classification	Supervised	UCI SpamBase
MNIST (Digits Classification)	Classification	Supervised	MNIST Dataset
Predicting Diabetes	Classification	Supervised	PIMA Indian Dataset
Iris Dataset	Classification	Supervised	UCI Iris Dataset

Dataset Summary with Feature Selection and Algorithms

Dataset	Type of ML Task	Feature Selection Technique	Suitable ML Algorithm
Iris Dataset	Classification	ANOVA (f.classif), SelectKBest	Logistic Regression, KNN
Loan Amount Prediction	Regression	SelectKBest (f_regression)	Linear Regression, Random Forest
Predicting Diabetes	Classification	Chi2, SelectKBest	SVM, Random Forest, XG-Boost
Email Spam Classification	Classification	Chi2, Mutual Info	Naive Bayes, Decision Tree
MNIST Handwritten Recognition	Classification	PCA or CNN-based feature selection	KNN, SVM, CNN (Deep Learning)

Learning Outcomes

- Gained practical experience in working with essential Python libraries such as **NumPy**, **Pandas**, **SciPy**, **Scikit-learn**, and **Matplotlib**.
- Understood the process of loading and preparing real-world datasets from public repositories like the UCI Machine Learning Repository and Kaggle.
- Learned to perform Exploratory Data Analysis (EDA) using summary statistics and visual tools such as histograms, heatmaps, and scatter plots.
- Applied data preprocessing techniques including handling missing values, encoding categorical variables, and feature scaling.
- Explored and applied feature selection techniques such as SelectKBest, chi2, and f.classif to improve model performance.
- Implemented end-to-end machine learning workflows for both classification and regression tasks.

Date: 23-07-2025

Experiment: 1

Name: Pranavah Varun

Roll No: 3122237001039

- Evaluated models using metrics like accuracy, confusion matrix, classification report, mean squared error, and R^2 score.
- Developed confidence in identifying suitable machine learning models based on the dataset and task type.