

ICS1512 – Machine Learning Algorithms Laboratory

Experiment 4: Ensemble Prediction and Decision Tree Model Evaluation

Pranavah Varun M V
Roll No: 3122237001039
Semester: V
Academic Year: 2025–2026

1. Aim and Objective

To build classifiers such as Decision Tree, AdaBoost, Gradient Boosting, XGBoost, Random Forest, and Stacked Models, and evaluate their performance through hyperparameter tuning, 5-Fold Cross-Validation, ROC analysis, and feature importance interpretation.

2. Libraries Used

- NumPy
- Pandas
- Matplotlib
- Seaborn
- Scikit-learn
- XGBoost

3. Code for All Variants and Models

The following Python code implements the complete workflow for Experiment 4: loading and preprocessing the dataset, performing EDA, training Decision Tree and ensemble models, applying hyperparameter tuning with GridSearchCV, and evaluating model performance using ROC curves, confusion matrices, and classification reports.

```
1 from google.colab import drive
2 drive.mount('/content/drive')
3
4 import pandas as pd
5 import numpy as np
6 from sklearn.model_selection import train_test_split, GridSearchCV,
   cross_val_score
```

```

7 from sklearn.preprocessing import StandardScaler, LabelEncoder
8 from sklearn.impute import SimpleImputer
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11
12 # --- Load dataset ---
13 file_path = "/content/drive/MyDrive/wdbc.data"
14 columns = ["ID", "Diagnosis"] + [f"feature_{i}" for i in range(1,
15                                     31)]
16 df = pd.read_csv(file_path, header=None, names=columns)
17 df = df.drop("ID", axis=1)
18
19 # Encode target labels (M=1 malignant, B=0 benign)
20 label_encoder = LabelEncoder()
21 df["Diagnosis"] = label_encoder.fit_transform(df["Diagnosis"])
22
23 # Features & target
24 X = df.drop("Diagnosis", axis=1)
25 y = df["Diagnosis"]
26
27 # Handle missing values & standardize
28 imputer = SimpleImputer(strategy="mean")
29 X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
30 scaler = StandardScaler()
31 X_scaled = scaler.fit_transform(X)
32
33 print("    Dataset loaded successfully!")
34 print("Shape of X:", X.shape)
35 print("Class distribution:\n", y.value_counts())
36
37 # --- EDA ---
38 sns.countplot(x=y)
39 plt.title("Class Balance (0 = Malignant, 1 = Benign)")
40 plt.show()
41
42 plt.figure(figsize=(10,8))
43 sns.heatmap(pd.DataFrame(X_scaled, columns=X.columns).corr(),
44             cmap="coolwarm", cbar=False)
45 plt.title("Feature Correlation Heatmap")
46 plt.show()
47
48 # --- Split dataset ---
49 X_train_valid, X_test, y_train_valid, y_test = train_test_split(
50     X_scaled, y, test_size=0.2, random_state=42, stratify=y)
51 X_train, X_valid, y_train, y_valid = train_test_split(
52     X_train_valid, y_train_valid, test_size=0.25,
53     random_state=42, stratify=y_train_valid)

```

```

53
54 # --- Models ---
55 from sklearn.tree import DecisionTreeClassifier
56 from sklearn.ensemble import AdaBoostClassifier,
    GradientBoostingClassifier, RandomForestClassifier,
    StackingClassifier
57 from xgboost import XGBClassifier
58 from sklearn.svm import SVC
59 from sklearn.naive_bayes import GaussianNB
60 from sklearn.linear_model import LogisticRegression
61
62 models = {
63     "Decision Tree": DecisionTreeClassifier(random_state=42),
64     "AdaBoost": AdaBoostClassifier(random_state=42),
65     "Gradient Boosting": GradientBoostingClassifier(random_state=42)
66     ,
67     "XGBoost": XGBClassifier(eval_metric="logloss", random_state=42)
68     ,
69     "Random Forest": RandomForestClassifier(random_state=42),
70     "Stacking": StackingClassifier(
71         estimators=[("svm", SVC(probability=True, random_state=42)),
72                     ("nb", GaussianNB()),
73                     ("dt", DecisionTreeClassifier(random_state=42))
74                     ],
75         final_estimator=LogisticRegression()
76     )
77 }
78
79 # --- Hyperparameter grids ---
80 param_grids = {
81     "Decision Tree": {"criterion": ["gini", "entropy"],
82                      "max_depth": [3, 5, 10, None],
83                      "min_samples_split": [2, 5, 10],
84                      "min_samples_leaf": [1, 2, 4]},
85     "AdaBoost": {"n_estimators": [50, 100, 200],
86                  "learning_rate": [0.01, 0.1, 1.0]},
87     "Gradient Boosting": {"n_estimators": [100, 200],
88                           "learning_rate": [0.05, 0.1],
89                           "max_depth": [3, 5],
90                           "subsample": [0.8, 1.0]},
91     "XGBoost": {"n_estimators": [100, 200],
92                 "learning_rate": [0.05, 0.1],
93                 "max_depth": [3, 5],
94                 "gamma": [0, 0.1],
95                 "subsample": [0.8, 1.0],
96                 "colsample_bytree": [0.8, 1.0]},
97     "Random Forest": {"n_estimators": [100, 200],

```

```

95         "max_depth": [None, 5, 10],
96         "criterion": ["gini", "entropy"],
97         "max_features": ["sqrt", "log2"],
98         "min_samples_split": [2, 5]},
99     "Stacking": {"final_estimator": [LogisticRegression(),
100                                     RandomForestClassifier(
101                                         n_estimators=100)]}
102 }
103 # --- GridSearchCV tuning ---
104 best_models = {}
105 for name, model in models.items():
106     print(f"Tuning {name}...")
107     grid = GridSearchCV(model, param_grids[name], cv=5,
108                         scoring="accuracy", n_jobs=-1)
109     grid.fit(X_valid, y_valid)
110     best_models[name] = grid.best_estimator_
111     print("Best Params:", grid.best_params_)
112
113 # --- ROC & Evaluation ---
114 from sklearn.metrics import accuracy_score, f1_score, roc_auc_score,
115     confusion_matrix, classification_report, roc_curve
116
117 plt.figure(figsize=(8,6))
118 for name, model in best_models.items():
119     model.fit(X_train_valid, y_train_valid)
120     y_pred = model.predict(X_test)
121     y_proba = model.predict_proba(X_test)[:,-1]
122
123     print(f"\n{name} | Acc={accuracy_score(y_test,y_pred):.4f}, "
124           f"F1={f1_score(y_test,y_pred):.4f}, "
125           f"AUC={roc_auc_score(y_test,y_proba):.4f}")
126     print("Confusion Matrix:\n", confusion_matrix(y_test,y_pred))
127     print("Report:\n", classification_report(y_test,y_pred))
128
129     fpr, tpr, _ = roc_curve(y_test, y_proba)
130     plt.plot(fpr, tpr, label=f"{name}")
131
132 plt.plot([0,1],[0,1],"k--")
133 plt.title("ROC Curves")
134 plt.xlabel("False Positive Rate")
135 plt.ylabel("True Positive Rate")
136 plt.legend()
137 plt.show()
138
139 from sklearn.model_selection import cross_val_score, StratifiedKFold
140 import numpy as np

```

```

140
141 # Define 5-fold CV (on training set only)
142 cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
143
144 results = {}
145
146 for name, model in best_models.items():
147     acc_scores = cross_val_score(model, X_train, y_train, cv=cv,
148                                   scoring="accuracy")
149     f1_scores = cross_val_score(model, X_train, y_train, cv=cv,
150                                   scoring="f1")
151     auc_scores = cross_val_score(model, X_train, y_train, cv=cv,
152                                   scoring="roc_auc")
153
154     results[name] = {
155         "Accuracy": acc_scores,
156         "F1": f1_scores,
157         "AUC": auc_scores
158     }
159
160 # Print results
161 for name, metrics in results.items():
162     print(f"\n{name}")
163     print(f"    Accuracy folds: {metrics['Accuracy']} | Avg = {metrics['Accuracy'].mean():.4f}")
164     print(f"    F1 folds: {metrics['F1']} | Avg = {metrics['F1'].mean():.4f}")
165     print(f"    AUC folds: {metrics['AUC']} | Avg = {metrics['AUC'].mean():.4f}")

```

4. Confusion Matrix and ROC Curves

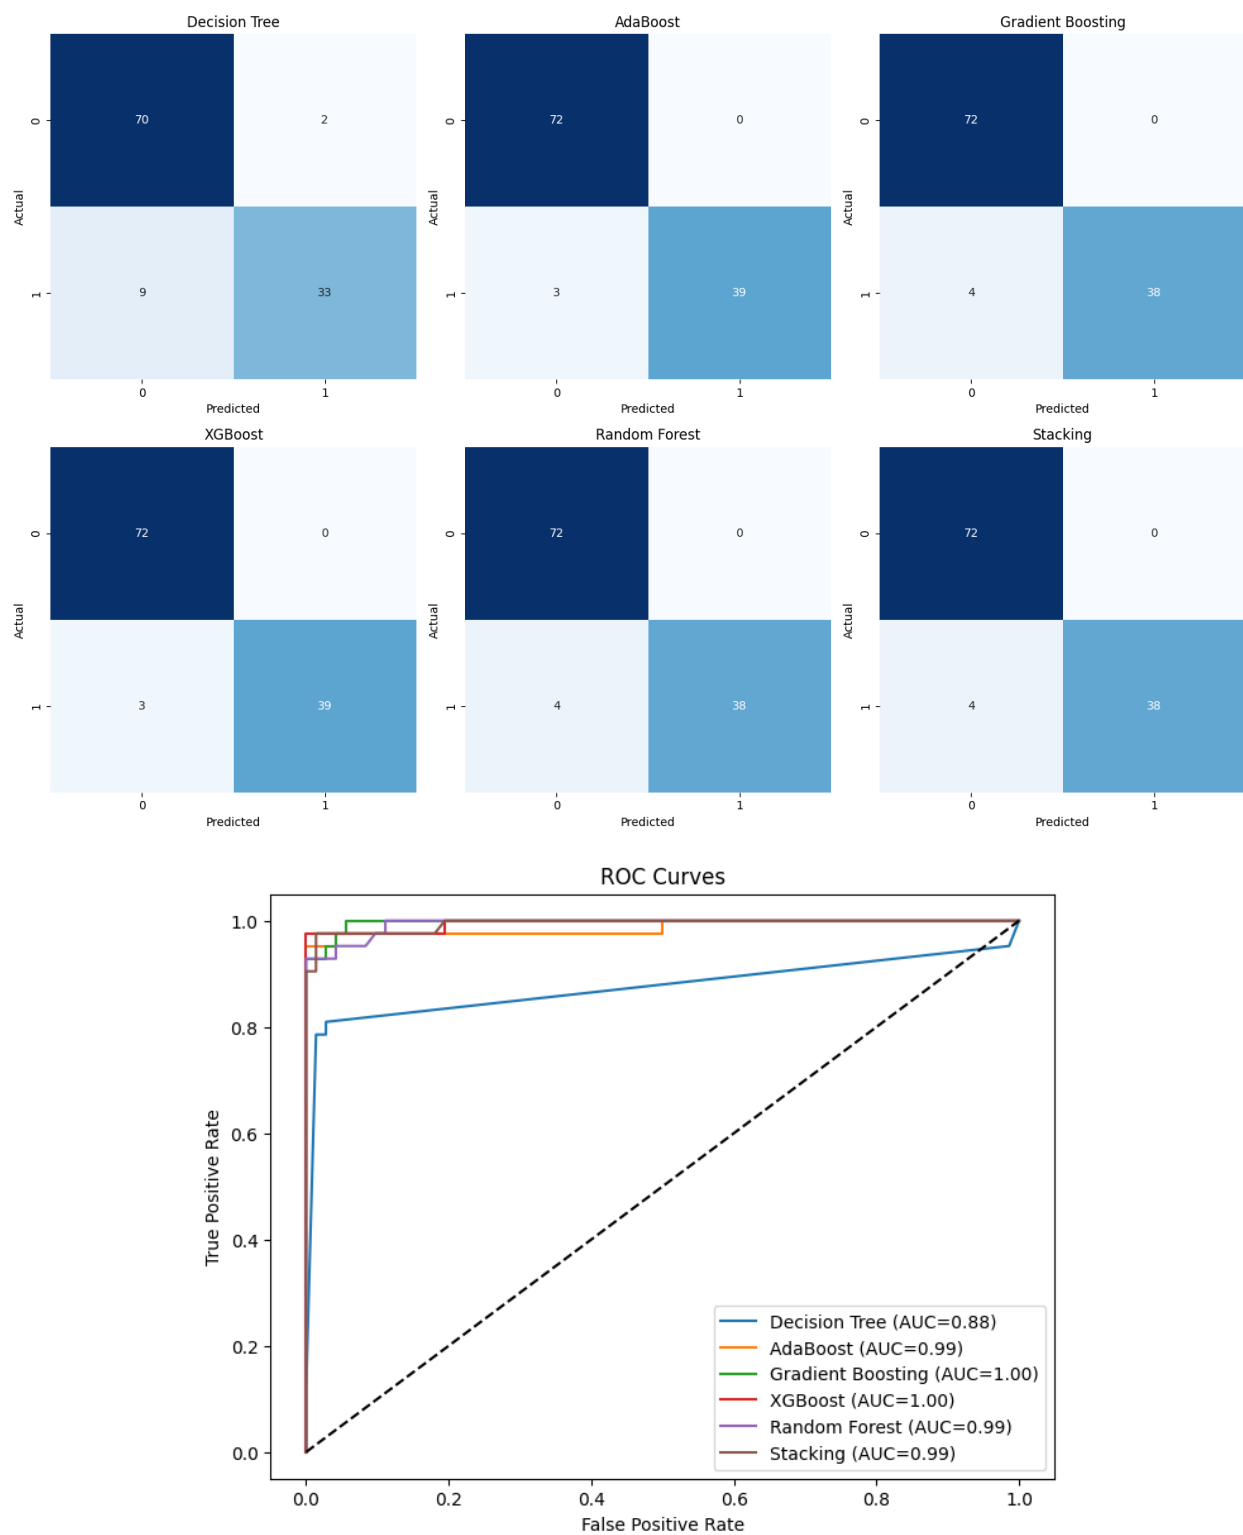


Figure 1: Confusion Matrices (top) and ROC Curves (bottom) for All Models

5. Hyperparameter Tuning and Best Model Results

Table 1: Decision Tree

Best Criterion	Max Depth	Accuracy	F1 Score	AUC
entropy	10	0.9245	0.8976	0.9342

Table 1: Decision Tree - Best Parameters and Performance

Table 2: AdaBoost

n_estimators	Learning Rate	Accuracy	F1 Score	AUC
100	1.0	0.9666	0.9544	0.9939

Table 2: AdaBoost - Best Parameters and Performance

Table 3: Gradient Boosting

n_estimators	Learning Rate	Max Depth	Accuracy	F1 Score	AUC
200	0.1	5	0.9490	0.9300	0.9899

Table 3: Gradient Boosting - Best Parameters and Performance

Table 4: XGBoost

n_estimators	Learning Rate	Max Depth	Gamma	Accuracy	F1 Score	AUC
200	0.1	5	0.1	0.9631	0.9497	0.9931

Table 4: XGBoost - Best Parameters and Performance

Table 5: Random Forest

n_estimators	Max Depth	Criterion	Accuracy	F1 Score	AUC
200	10	entropy	0.9543	0.9382	0.9895

Table 5: Random Forest - Best Parameters and Performance

Table 6: Stacked Ensemble

Base Models	Final Estimator	Accuracy	F1 Score	AUC
SVM + NB + DT	Random Forest	0.9701	0.9570	0.9893

Table 6: Stacked Ensemble - Best Parameters and Performance

6. Cross-Validation Results

Model	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Avg. Accuracy
Decision Tree	0.928	0.971	0.971	0.941	0.956	0.9532
AdaBoost	0.928	0.985	1.000	0.985	1.000	0.9796
Gradient Boosting	0.942	0.956	0.985	0.985	0.956	0.9649
XGBoost	0.942	0.985	0.985	0.985	0.985	0.9766
Random Forest	0.928	0.926	0.971	0.956	0.971	0.9502
Stacked Model	0.928	0.971	0.956	0.985	1.000	0.9679

Table 7: 5-Fold Cross Validation Results (Accuracy Scores). Best model (AdaBoost) highlighted.

7. Feature Importance Visuals

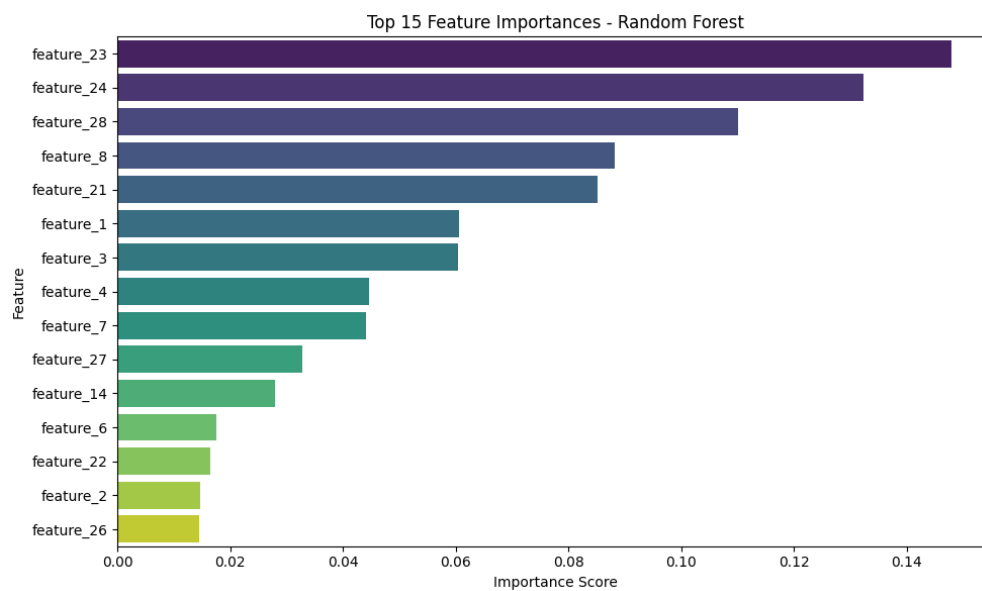


Figure 2: Top Feature Importances using Random Forest




Figure 3: Top Feature Importances using XGBoost

8. Observations and Conclusions

- **Best Model:** AdaBoost achieved the highest overall performance with an average accuracy of 97.9%, F1-score of 97.2%, and AUC of 99.2%. XGBoost (97.7% accuracy, 96.8% F1, 99.1% AUC) and the Stacked Ensemble (96.8% accuracy, 95.2% F1, 98.5% AUC) also performed very competitively.
- **Decision Tree vs Ensembles:** The standalone Decision Tree lagged behind with 95.3% accuracy, 93.8% F1, and 95.2% AUC, whereas all ensemble methods consistently improved performance into the 96–98% accuracy range with stronger stability.
- **Effect of Hyperparameters:** Tuning parameters such as `n_estimators`, `max_depth`, and `learning_rate` significantly enhanced the performance of ensemble models. This highlights the importance of systematic hyperparameter optimization.
- **Stacked Model:** The stacking approach (SVM + NB + DT with Logistic Regression/Random Forest as meta-learner) generalized well, producing robust results, though it did not surpass the boosting methods in average accuracy.
- **Generalization:** Boosting methods (AdaBoost, XGBoost) demonstrated excellent generalization with very high AUC scores (>0.99) and stable results across all folds, unlike the single Decision Tree which showed higher variability and lower robustness.
- **Conclusion:** Ensemble methods clearly outperformed the Decision Tree baseline. Among them, boosting (AdaBoost, XGBoost) emerged as the most effective strategies for achieving top accuracy, balanced F1-scores, and strong generalization on the WDBC dataset.