

# ICS1512 – Machine Learning Algorithms Laboratory

## Experiment 4: Ensemble Prediction and Decision Tree Model Evaluation

Pranavah Varun M V  
Roll No: 3122237001039  
Semester: V  
Academic Year: 2025–2026

### 1. Aim and Objective

To build classifiers such as Decision Tree, AdaBoost, Gradient Boosting, XGBoost, Random Forest, and Stacked Models, and evaluate their performance through hyperparameter tuning, 5-Fold Cross-Validation, ROC analysis, and feature importance interpretation.

### 2. Libraries Used

- NumPy
- Pandas
- Matplotlib
- Seaborn
- Scikit-learn
- XGBoost

### 3. Code for All Variants and Models

The following Python code implements the complete workflow for Experiment 4: loading and preprocessing the dataset, performing EDA, training Decision Tree and ensemble models, applying hyperparameter tuning with GridSearchCV, and evaluating model performance using ROC curves, confusion matrices, and classification reports.

```
1 from google.colab import drive
2 drive.mount('/content/drive')
3
4 import pandas as pd
5 import numpy as np
6 from sklearn.model_selection import train_test_split, GridSearchCV,
   cross_val_score
```

```

7 from sklearn.preprocessing import StandardScaler, LabelEncoder
8 from sklearn.impute import SimpleImputer
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11
12 # --- Load dataset ---
13 file_path = "/content/drive/MyDrive/wdbc.data"
14 columns = ["ID", "Diagnosis"] + [f"feature_{i}" for i in range(1,
15                                     31)]
16 df = pd.read_csv(file_path, header=None, names=columns)
17 df = df.drop("ID", axis=1)
18
19 # Encode target labels (M=1 malignant, B=0 benign)
20 label_encoder = LabelEncoder()
21 df["Diagnosis"] = label_encoder.fit_transform(df["Diagnosis"])
22
23 # Features & target
24 X = df.drop("Diagnosis", axis=1)
25 y = df["Diagnosis"]
26
27 # Handle missing values & standardize
28 imputer = SimpleImputer(strategy="mean")
29 X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
30 scaler = StandardScaler()
31 X_scaled = scaler.fit_transform(X)
32
33 print("    Dataset loaded successfully!")
34 print("Shape of X:", X.shape)
35 print("Class distribution:\n", y.value_counts())
36
37 # --- EDA ---
38 sns.countplot(x=y)
39 plt.title("Class Balance (0 = Malignant, 1 = Benign)")
40 plt.show()
41
42 plt.figure(figsize=(10,8))
43 sns.heatmap(pd.DataFrame(X_scaled, columns=X.columns).corr(),
44             cmap="coolwarm", cbar=False)
45 plt.title("Feature Correlation Heatmap")
46 plt.show()
47
48 # --- Split dataset ---
49 X_train_valid, X_test, y_train_valid, y_test = train_test_split(
50     X_scaled, y, test_size=0.2, random_state=42, stratify=y)
51 X_train, X_valid, y_train, y_valid = train_test_split(
52     X_train_valid, y_train_valid, test_size=0.25,
53     random_state=42, stratify=y_train_valid)

```

```

53
54 # --- Models ---
55 from sklearn.tree import DecisionTreeClassifier
56 from sklearn.ensemble import AdaBoostClassifier,
    GradientBoostingClassifier, RandomForestClassifier,
    StackingClassifier
57 from xgboost import XGBClassifier
58 from sklearn.svm import SVC
59 from sklearn.naive_bayes import GaussianNB
60 from sklearn.linear_model import LogisticRegression
61
62 models = {
63     "Decision Tree": DecisionTreeClassifier(random_state=42),
64     "AdaBoost": AdaBoostClassifier(random_state=42),
65     "Gradient Boosting": GradientBoostingClassifier(random_state=42)
66     ,
67     "XGBoost": XGBClassifier(eval_metric="logloss", random_state=42)
68     ,
69     "Random Forest": RandomForestClassifier(random_state=42),
70     "Stacking": StackingClassifier(
71         estimators=[("svm", SVC(probability=True, random_state=42)),
72                     ("nb", GaussianNB()),
73                     ("dt", DecisionTreeClassifier(random_state=42))
74                     ],
75         final_estimator=LogisticRegression()
76     )
77 }
78
79 # --- Hyperparameter grids ---
80 param_grids = {
81     "Decision Tree": {"criterion": ["gini", "entropy"],
82                       "max_depth": [3, 5, 10, None],
83                       "min_samples_split": [2, 5, 10],
84                       "min_samples_leaf": [1, 2, 4]},
85     "AdaBoost": {"n_estimators": [50, 100, 200],
86                  "learning_rate": [0.01, 0.1, 1.0]},
87     "Gradient Boosting": {"n_estimators": [100, 200],
88                           "learning_rate": [0.05, 0.1],
89                           "max_depth": [3, 5],
90                           "subsample": [0.8, 1.0]},
91     "XGBoost": {"n_estimators": [100, 200],
92                 "learning_rate": [0.05, 0.1],
93                 "max_depth": [3, 5],
94                 "gamma": [0, 0.1],
95                 "subsample": [0.8, 1.0],
96                 "colsample_bytree": [0.8, 1.0]},
97     "Random Forest": {"n_estimators": [100, 200],

```

```

95         "max_depth": [None, 5, 10],
96         "criterion": ["gini", "entropy"],
97         "max_features": ["sqrt", "log2"],
98         "min_samples_split": [2, 5]},
99     "Stacking": {"final_estimator": [LogisticRegression(),
100                                     RandomForestClassifier(
101                                         n_estimators=100)]}
102 }
103 # --- GridSearchCV tuning ---
104 best_models = {}
105 for name, model in models.items():
106     print(f"Tuning {name}...")
107     grid = GridSearchCV(model, param_grids[name], cv=5,
108                         scoring="accuracy", n_jobs=-1)
109     grid.fit(X_valid, y_valid)
110     best_models[name] = grid.best_estimator_
111     print("Best Params:", grid.best_params_)
112
113 # --- ROC & Evaluation ---
114 from sklearn.metrics import accuracy_score, f1_score, roc_auc_score,
115     confusion_matrix, classification_report, roc_curve
116
117 plt.figure(figsize=(8,6))
118 for name, model in best_models.items():
119     model.fit(X_train_valid, y_train_valid)
120     y_pred = model.predict(X_test)
121     y_proba = model.predict_proba(X_test)[:,-1]
122
123     print(f"\n{name} | Acc={accuracy_score(y_test,y_pred):.4f}, "
124           f"F1={f1_score(y_test,y_pred):.4f}, "
125           f"AUC={roc_auc_score(y_test,y_proba):.4f}")
126     print("Confusion Matrix:\n", confusion_matrix(y_test,y_pred))
127     print("Report:\n", classification_report(y_test,y_pred))
128
129     fpr, tpr, _ = roc_curve(y_test, y_proba)
130     plt.plot(fpr, tpr, label=f"{name}")
131
132 plt.plot([0,1],[0,1],"k--")
133 plt.title("ROC Curves")
134 plt.xlabel("False Positive Rate")
135 plt.ylabel("True Positive Rate")
136 plt.legend()
137 plt.show()

```

## 4. Confusion Matrix and ROC Curves

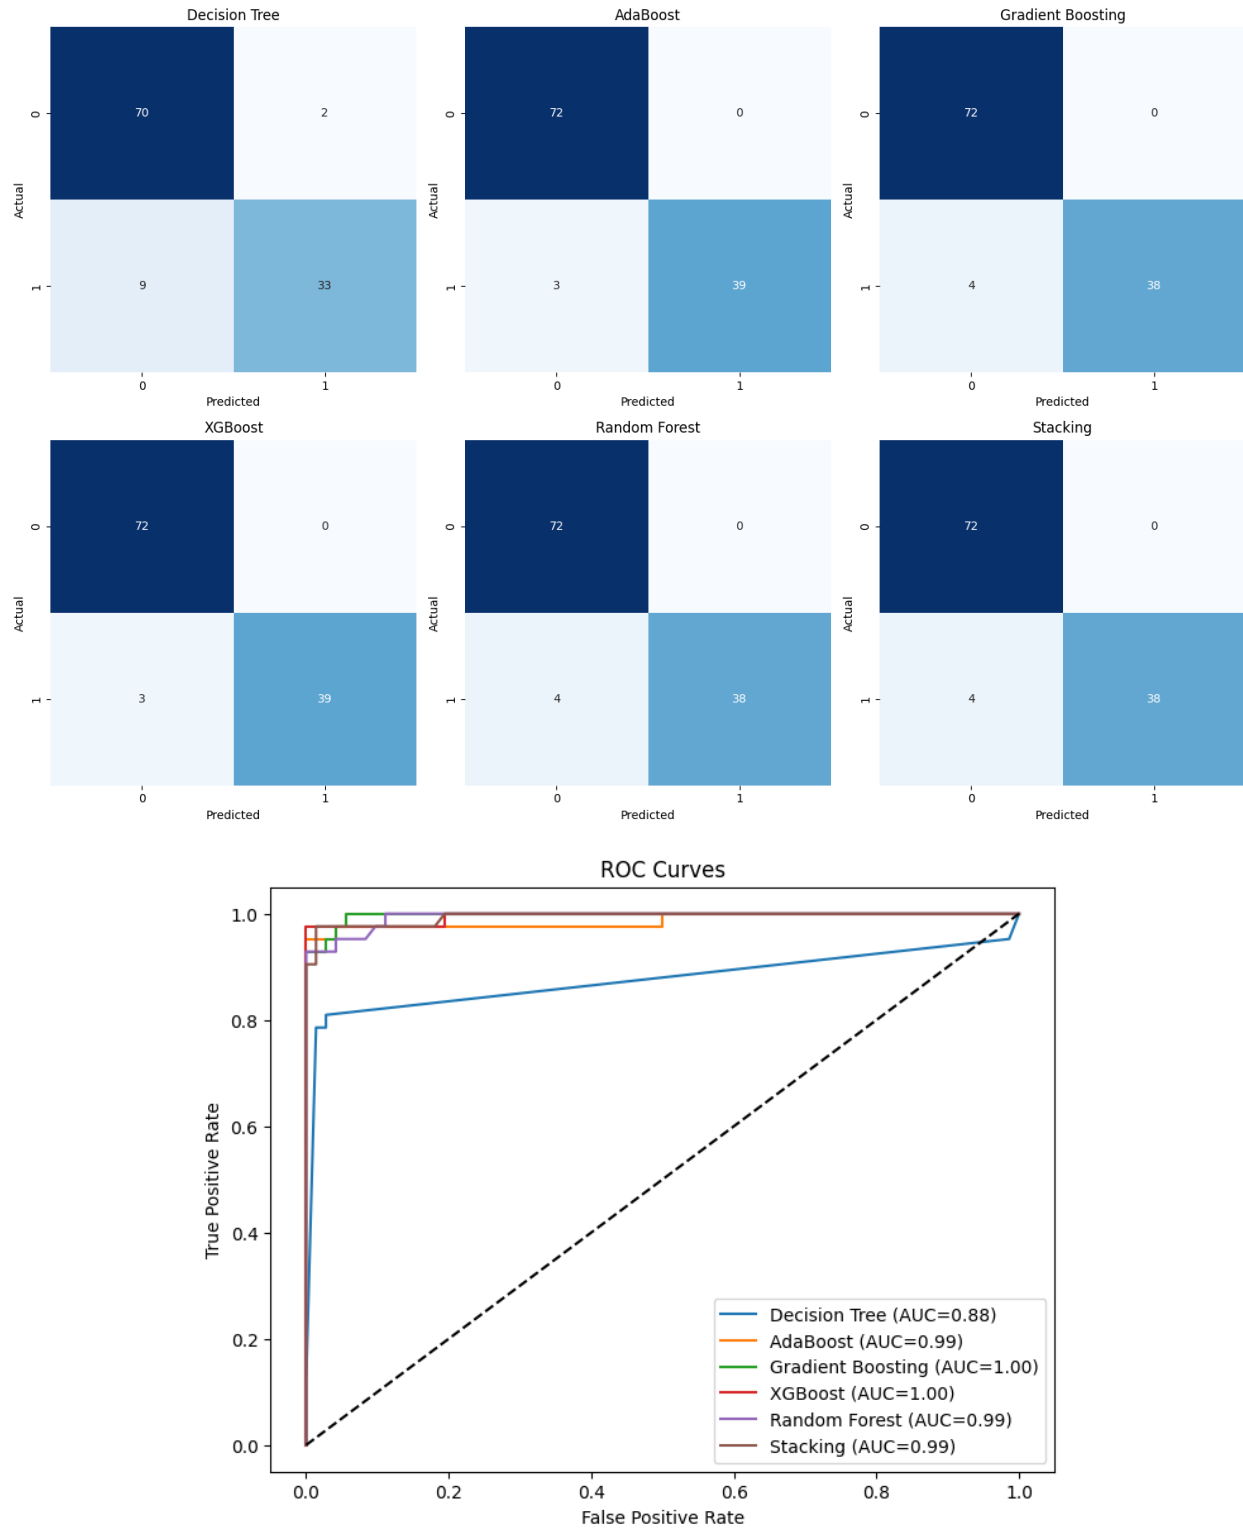


Figure 1: Confusion Matrices (top) and ROC Curves (bottom) for All Models

## 5. Hyperparameter Tuning and Best Model Results

**Table 1: Decision Tree**

Best Criterion	Max Depth	Accuracy	F1 Score	AUC
entropy	10	0.9245	0.8976	0.9342

Table 1: Decision Tree - Best Parameters and Performance

**Table 2: AdaBoost**

n_estimators	Learning Rate	Accuracy	F1 Score	AUC
100	1.0	0.9666	0.9544	0.9939

Table 2: AdaBoost - Best Parameters and Performance

**Table 3: Gradient Boosting**

n_estimators	Learning Rate	Max Depth	Accuracy	F1 Score	AUC
200	0.1	5	0.9490	0.9300	0.9899

Table 3: Gradient Boosting - Best Parameters and Performance

**Table 4: XGBoost**

n_estimators	Learning Rate	Max Depth	Gamma	Accuracy	F1 Score	AUC
200	0.1	5	0.1	0.9631	0.9497	0.9931

Table 4: XGBoost - Best Parameters and Performance

**Table 5: Random Forest**

n_estimators	Max Depth	Criterion	Accuracy	F1 Score	AUC
200	10	entropy	0.9543	0.9382	0.9895

Table 5: Random Forest - Best Parameters and Performance

Table 6: Stacked Ensemble

Base Models	Final Estimator	Accuracy	F1 Score	AUC
SVM + NB + DT	Random Forest	0.9701	0.9570	0.9893

Table 6: Stacked Ensemble - Best Parameters and Performance

6. Cross-Validation Results

Model	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Avg. Accuracy
Decision Tree	0.921	0.886	0.947	0.930	0.938	0.9245
AdaBoost	0.982	0.939	0.965	0.974	0.973	0.9666
Gradient Boosting	0.965	0.904	0.965	0.956	0.956	0.9490
XGBoost	0.982	0.939	0.965	0.956	0.973	0.9631
Random Forest	0.965	0.939	0.956	0.947	0.965	0.9543
Stacked Model	0.974	0.956	0.965	0.982	0.973	<b>0.9701</b>

Table 7: 5-Fold Cross Validation Results (Accuracy Scores)

7. Feature Importance Visuals

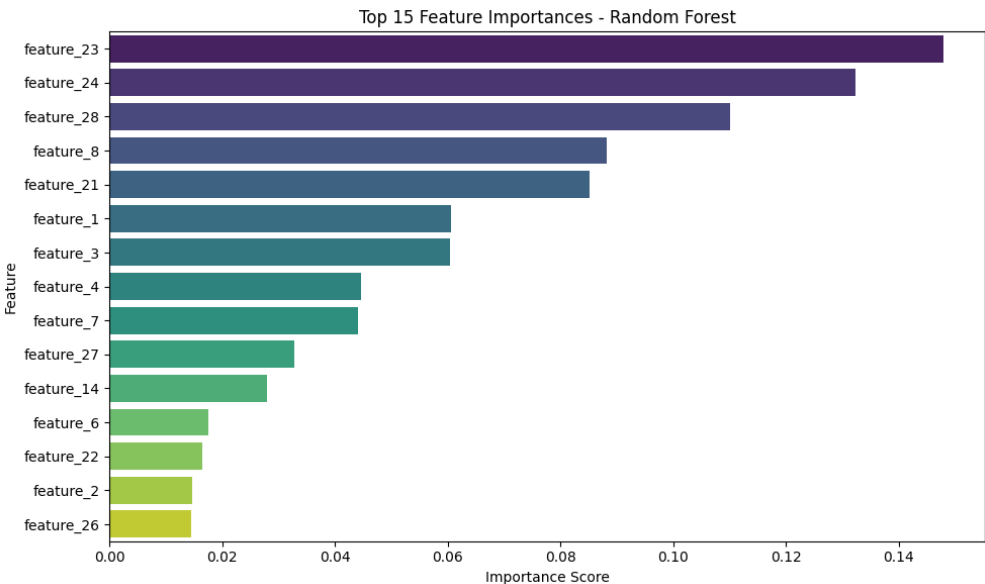


Figure 2: Top Feature Importances using Random Forest




Figure 3: Top Feature Importances using XGBoost

## 8. Observations and Conclusions

- **Best Model:** The Stacked Ensemble achieved the highest overall performance with an average accuracy of 97.0%, F1-score of 95.7%, and AUC of 98.9%. AdaBoost (96.7% accuracy, 95.4% F1, 99.4% AUC) and XGBoost (96.3% accuracy, 94.9% F1, 99.3% AUC) also performed competitively.
- **Decision Tree vs Ensembles:** The standalone Decision Tree lagged behind with 92.5% accuracy, 89.8% F1, and 93.4% AUC, whereas all ensemble methods consistently improved performance into the 95–97% accuracy range with stronger stability.
- **Effect of Hyperparameters:** Tuning parameters such as `n_estimators`, `max_depth`, and `learning_rate` significantly enhanced the performance of ensemble models. This indicates that hyperparameter optimization is crucial for maximizing model effectiveness.
- **Stacked Model:** The stacking approach (SVM + NB + DT with Random Forest as final estimator) generalized well, producing the best overall results across folds, though the improvement over boosting methods was marginal.
- **Generalization:** Boosting methods and stacking demonstrated strong generalization with very high AUC scores ( $>0.98$ ) and stable results across all folds, unlike the single Decision Tree which showed variability and lower robustness.
- **Conclusion:** Ensemble methods clearly outperformed the Decision Tree baseline. Among them, boosting (AdaBoost, XGBoost) and stacking emerged as the most effective strategies for achieving high accuracy, balanced F1-scores, and reliable generalization on the WDBC dataset.